

S4: Time Series Analysis, Time Series Preprocessing and Modeling with ML

Phanuphat Srisukhawasu (Oad)

Machine Learning Department
NUS Fintech Society

Week 6 (16 Sep 2024 - 20 Sep 2024)



Session Outline

Introduction to Time Series Analysis

- What are Time Series?

- Characteristics of Time Series

Time Series Preprocessing

- Setting Dataframe Index to Date and Time

- Handling Missing Time Series Data

- Aggregating Time Series Data

- Creating Time Series Features

Time Series Modeling

- Introduction to ARIMA

- Time Series Forecasting with ARIMA

Machine Learning Models for Time Series

- Splitting Time Series Data


- Using Random Forests and XGBoost

- Evaluating Machine Learning Model Performance


- Time Series Forecasting in Scikit-learn Cheat Sheet

Time Series Data in Simple Words

- ▶ **Time series** are data points indexed in an order of time.
- ▶ It could be sampled as frequent as we want: **every second, every minute, hourly, daily, monthly, yearly, etc.**



Trend, Seasonality, and Noise



Note: **Residual** (a superset of noises) is **what we are left with**, after analyzing the other components.

Loading Time Series Data

- ▶ As we have a CSV file containing time series data, we could use `pd.read_csv()` to import your data as usual.
- ▶ But when we inspect the data, we generally see that the **data type of the date and time column is a string (object.)**

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1226 entries, 0 to 1225
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Date            1226 non-null   object
1   Open            1226 non-null   int64
2   High            1226 non-null   int64
3   Low             1226 non-null   int64
4   Close           1226 non-null   int64
5   Volume          1226 non-null   int64
6   Stock Trading   1226 non-null   int64
dtypes: int64(6), object(1)
memory usage: 67.2+ KB
```

Converting Date Column to Datetime Object

- ▶ In general, we will process the dataframe in such way that the Date column is a Datetime object instead of a string.
- ▶ To accomplish this, we usually use `df['Date'] = pd.to_datetime(df['Date'])`.
- ▶ We also usually do `df = df.set_index('Date')`. This will make the processing more convenient.


	Open	High	Low	Close	Volume	Stock Trading
Date						
2016-12-30	42120	42330	41700	41830	610000	25628028000
2016-12-29	43000	43220	42540	42660	448400	19188227000
2016-12-28	43940	43970	43270	43270	339900	14780670000
2016-12-27	43140	43700	43140	43620	400100	17427993000
2016-12-26	43310	43660	43090	43340	358200	15547803000


```

<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1226 entries, 2016-12-30 to 2012-01-04
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Open            1226 non-null   int64
1   High            1226 non-null   int64
2   Low             1226 non-null   int64
3   Close           1226 non-null   int64
4   Volume          1226 non-null   int64
5   Stock Trading   1226 non-null   int64
dtypes: int64(6)
memory usage: 67.0 KB
    
```

Filling the Missing Data with Forward Fill

- ▶ For time series, a simple method to fill missing values at time t_0 is to use the value from the closest previous time $t < t_0$.
- ▶ For this example dataset, we already know that we don't have any missing data. But if we do, we could use `df = df.ffill()`. After that, we could use `sns.lineplot()` from Seaborn to visualize the data.



Time Series Downsampling

- ▶ In this example, we have a data recorded in a **daily basis**. In some cases, we may want to analyze the data in a **weekly, monthly, or yearly basis**.
- ▶ In Pandas, we could resample the data as frequent as we want: `df = df.resample(<FREQ>).mean()`.



Note: The above plot is generated from sampling the data using 'W' as <FREQ> and calculate the average. This will resample the data **weekly**. The plot is more smoothed out compared to the previous. For other frequencies, you may refer to this [link](#).

Lag Features (1)

- ▶ To use machine learning, we need inputs or features. What are the **features for time series**?
- ▶ It does not make any senses to use the time index as the features, as the **date should not be correlated to the values**.
- ▶ In time series forecasting, we want to use values from the **past to predict the future** values.
- ▶ Using this principle, we can create features from the past values known as **lag features**.
- ▶ For example, we can use `df['lag1'] = df['Close'].shift(1)`. This command creates a new column 'lag1' that shows the closing price from one time step earlier.

Lag Features (2)

- ▶ In general, we can pass in different values of steps other than 1 (i.e. 2, 3, ...) to create lagging from multiple time steps.
- ▶ This allow us to have **multiple features** to train the model.

	Open	High	Low	Close	Volume	Stock Trading	lag1	lag2	lag3
Date									
2012-01-08	13920.0	13973.333333	13696.666667	13790.0	6.120333e+05	8.462075e+09	NaN	NaN	NaN
2012-01-15	14377.5	14630.000000	14230.000000	14567.5	1.019975e+06	1.480215e+10	13790.0	NaN	NaN
2012-01-22	14800.0	14908.000000	14702.000000	14794.0	5.475800e+05	8.120263e+09	14567.5	13790.0	NaN
2012-01-29	14832.0	15010.000000	14772.000000	14948.0	5.902400e+05	8.795512e+09	14794.0	14567.5	13790.0
2012-02-05	15212.0	15330.000000	15090.000000	15256.0	5.479400e+05	8.341651e+09	14948.0	14794.0	14567.5
2012-02-12	15742.0	15846.000000	15586.000000	15722.0	6.961400e+05	1.095340e+10	15256.0	14948.0	14794.0
2012-02-19	15826.0	16064.000000	15718.000000	15974.0	7.001000e+05	1.121111e+10	15722.0	15256.0	14948.0

Note: The **numbers of lagging features** to use could be one of the **hyperparameters** to tune.

ARIMA: AR (Autoregressive) + I (Integrated) + MA (Moving Average)


- ▶ **ARIMA** is a **statistical model** used to forecast the time series based on the past values.
- ▶ Sometimes, we could solve the simpler problems using just ARIMA and **without any machine learning!**
- ▶ In simple words, each part of ARIMA play the following roles:
 - ▶ **AR** = Previous values (**lag features**,)
 - ▶ **I** = Adjustments to make data stable (**removing trends**,)
 - ▶ **MA** = Adjustments to **smooth out the noises**.
- ▶ If you want to explore the mathematics behind ARIMA, feel free to visit this site (click on the formula XD):

$$y_t = \phi_1 y_{t-1} + \phi_2 y_{t-2} + \cdots + \phi_p y_{t-p} + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \cdots + \theta_q \epsilon_{t-q} + \epsilon_t$$

Implementing ARIMA in Python

- ▶ **ARIMA** is in a library called `statsmodels`. We can import it using `from statsmodels.tsa.arima.model import ARIMA`.
- ▶ We can **fit the model** to the data using: `model = ARIMA(df['Close'], order=(p, d, q)).fit()`.
 - ▶ `p` is the **number of lag features** to use (i.e. `p = 2` will use only the values at step $t - 1$ and $t - 2$.)
 - ▶ `d` is the **number of times we differenced the time series** ($x(t) - x(t - 1)$,) to make it stationary. We start with $d = 0$.
 - ▶ `q` is the **number of lagged forecast errors** to use.
- ▶ After fitting the model, we could **forecast** the time series values by assigning `predictions = model.predict(start=0, end=len(df)-1)`.


Example Results of ARIMA



Note: Sometimes, it is difficult to judge the **performance** by eyes. In the next section, we will introduce **numerical metrics to evaluate the models**.

Train-Test Split on Time Series

- ▶ When it comes to time series, we **cannot use random subsets** of data as the **training and testing sets**.
- ▶ Obviously, we want to train the model on **the past data to forecast the future**. Therefore, this is what we usually do:




More Advanced Machine Learning Models

- ▶ Last week, you have already learned about **Decision Trees**. In this session, we are introducing more advanced tree-based models which are **constructed by ensembling multiple trees**.
- ▶ With that being said, we shall introduce two primary methods for ensembling multiple trees:
 - ▶ **Bagging**: We construct multiple decision trees using **random subsets of data and features**. The final model decision is the **voting (average prediction)** from all of the trees. The model that use this ensembling method is **Random Forest**.
 - ▶ **Boosting**: We construct **new trees from the previous trees**. The early trees are called the **weak learner**. The errors made by them are taken into account as **the weights of importance** to train the subsequent trees, creating **strong learner** at the end. The models that use this ensembling method are **XGBoost, LightGBM**, and their variations.


Bagging and Boosting Illustration

Bagging



Parallel

Boosting



Sequential

Note: you can read more about the mathematical details of each model [here](#).

Implementing Random Forests & XGBoost in Python

- ▶ Like most models, Random Forests can also be accessed through Scikit-learn library. We can import it using `from sklearn.ensemble import RandomForestClassifier` or `RandomForestRegressor`.
- ▶ However, XGBoost is **not** inside Scikit-learn. We can import it from the library `xgboost` where we could use `XGBClassifier` and `XGBRegressor`. You may need to install it first if you haven't (`pip install xgboost`.)
- ▶ After having the models, you could use `model.fit()` on the training data as well as the other methods you have learned.
- ▶ Now, how do we evaluate the models using **metrics**?

Evaluating Time Series Forecasting Model

- ▶ In time series forecasting, we are trying to predict the **numerical outputs based on inputs**. It's just **regression!**
- ▶ To evaluate regression models, we usually look at the **average errors** they made (lower is better.)
 - ▶ **Root Mean Squared Error (RMSE):**

$$RMSE = \sqrt{\frac{\sum_{i=1}^N (y_i - \hat{y}_i)^2}{N}}$$

- ▶ **Mean Absolute Error (MAE):**

$$MAE = \frac{\sum_{i=1}^N |y_i - \hat{y}_i|}{N}$$

Note: Sometimes, we may also want to look at the **coefficient of determination (r2_score.)** This value range from 0 – 1, **higher is better.**

Data Preprocessing Summary

Here are the **general steps** to process time series data.

1. Convert the date and time column from string to Pandas' Datetime object using `df['Date'] = pd.to_datetime(df['Date'], format=<...>).`
2. Set the dataframe index to be that date and time column using `df = df.set_index('Date').`
3. Deal with the missing data with the appropriate methods. You could try a forward fill using `df = df.ffill().`
4. Resample the data to be indexed at the frequency you need (i.e. weekly, daily, hourly, etc.) using `df = df.resample(<FREQ>).mean().`
5. Generate multiple lag features and assigning them to multiple columns using `df['<LAG>'] = df['Close'].shift(<STEP>).`

Model Training Summary

1. Split the data into **training and testing** by simply **slicing the dataframe**. If we have enough data, we could split the data into half training and half testing (i.e. `df_train = df.loc[:len(df)//2, :]` and `df_test = df.loc[len(df)//2:, :]`.)
2. Create **features and target** (for each training and testing):
`X_train = df_train.loc[:, ['lag1', 'lag2', 'lag3', ...]]` and `y_train = df_train.loc[:, 'Close']`.
3. Train the model on the training data. For example:
 - ▶ Using **Random Forests**: `model = RandomForestRegressor()` and then `model.fit(X_train, y_train)`.
 - ▶ Using **XGBoost**: `model = XGBRegressor()` and then `model.fit(X_train, y_train)`.

Model Evaluation Summary

1. View the predicted time series by using `y_train_pred = model.predict(X_train)` and/or `y_test_pred = model.predict(X_test)`, then plot the results.
2. Evaluate the numerical metrics (**RMSE, MAE, and R-Squared**) using functions from `sklearn.metrics`. For example:
 - ▶ `mae = mean_absolute_error(y_test, y_pred_test)`
 - ▶ `rmse = root_mean_squared_error(y_test, y_pred_test)`
 - ▶ `r2 = r2_score(y_test, y_pred_test)`


Recap: Generally, we want to evaluate the model performance on the **testing set**, which reflects how well can the **model generalize to unseen data**.
3. Using the performance metrics, we can iteratively go back and **tune the models until getting the desired performance**.

Extra: Time Series Cross-Validation (1)

1. We know from the previous sessions that we use **cross-validation** to properly tune the **hyperparameters** while having a separated test set as the **actual unseen data**.
2. For time series, the appropriate way for cross validation is to use `sklearn.model_selection.TimeSeriesSplit()`
3. You may also pass an argument `gap` which determines how many samples you want to **exclude** between the training and validation set.
4. After assigning `cv = TimeSeriesSplit(n=5, gap=1)`, we can pass `cv` to classes like `GridSearchCV(cv=cv)` or `RandomizedSearchCV(cv=cv)` and then tune the hyperparameters as you wished.

Extra: Time Series Cross-Validation (2)

- ▶ This is how the cross-validation looks like when using `TimeSeriesSplit(n_splits=4)`.



- ▶ Note: More complex models like **Random Forests** and **XGBoost** have lots of hyperparameters to tune. You may consider looking at a **lightweight version of XGBoost**, **LightGBM** as well.